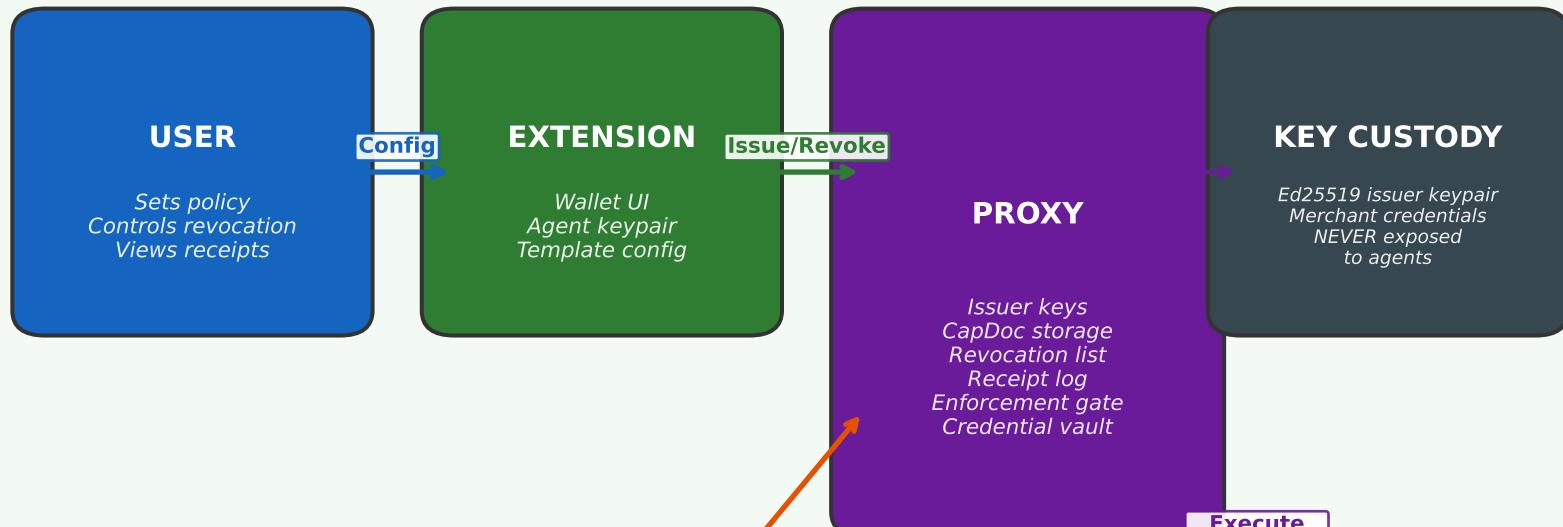


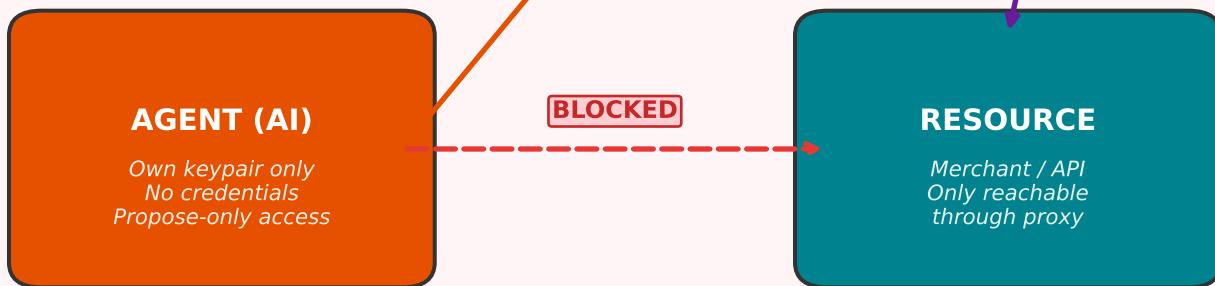
CAPNET SYSTEM ARCHITECTURE

Trust Boundaries & Component Roles

TRUSTED ZONE



UNTRUSTED ZONE



CAPABILITY ISSUANCE FLOW

How a capability is created and bound to an agent

USER

EXTENSION

PROXY

AGENT

User sets policy template:
"\$200, groceries, no alcohol, 7 days"

1

Extension sends
POST /capability/issue
with policy + agent pubkey

2

Proxy creates CapDoc:
• Generates cap_id
• Sets constraints from policy
• Binds to agent pubkey
• Signs with issuer key
• Stores locally
• Emits CAP_ISSUED receipt

3

Returns signed CapDoc

4

Shows "Capability Active"
with details + revoke button

5

Agent knows:
✓ A capability exists for it
✓ Its own keypair
✗ Merchant credentials
✗ Issuer signing key

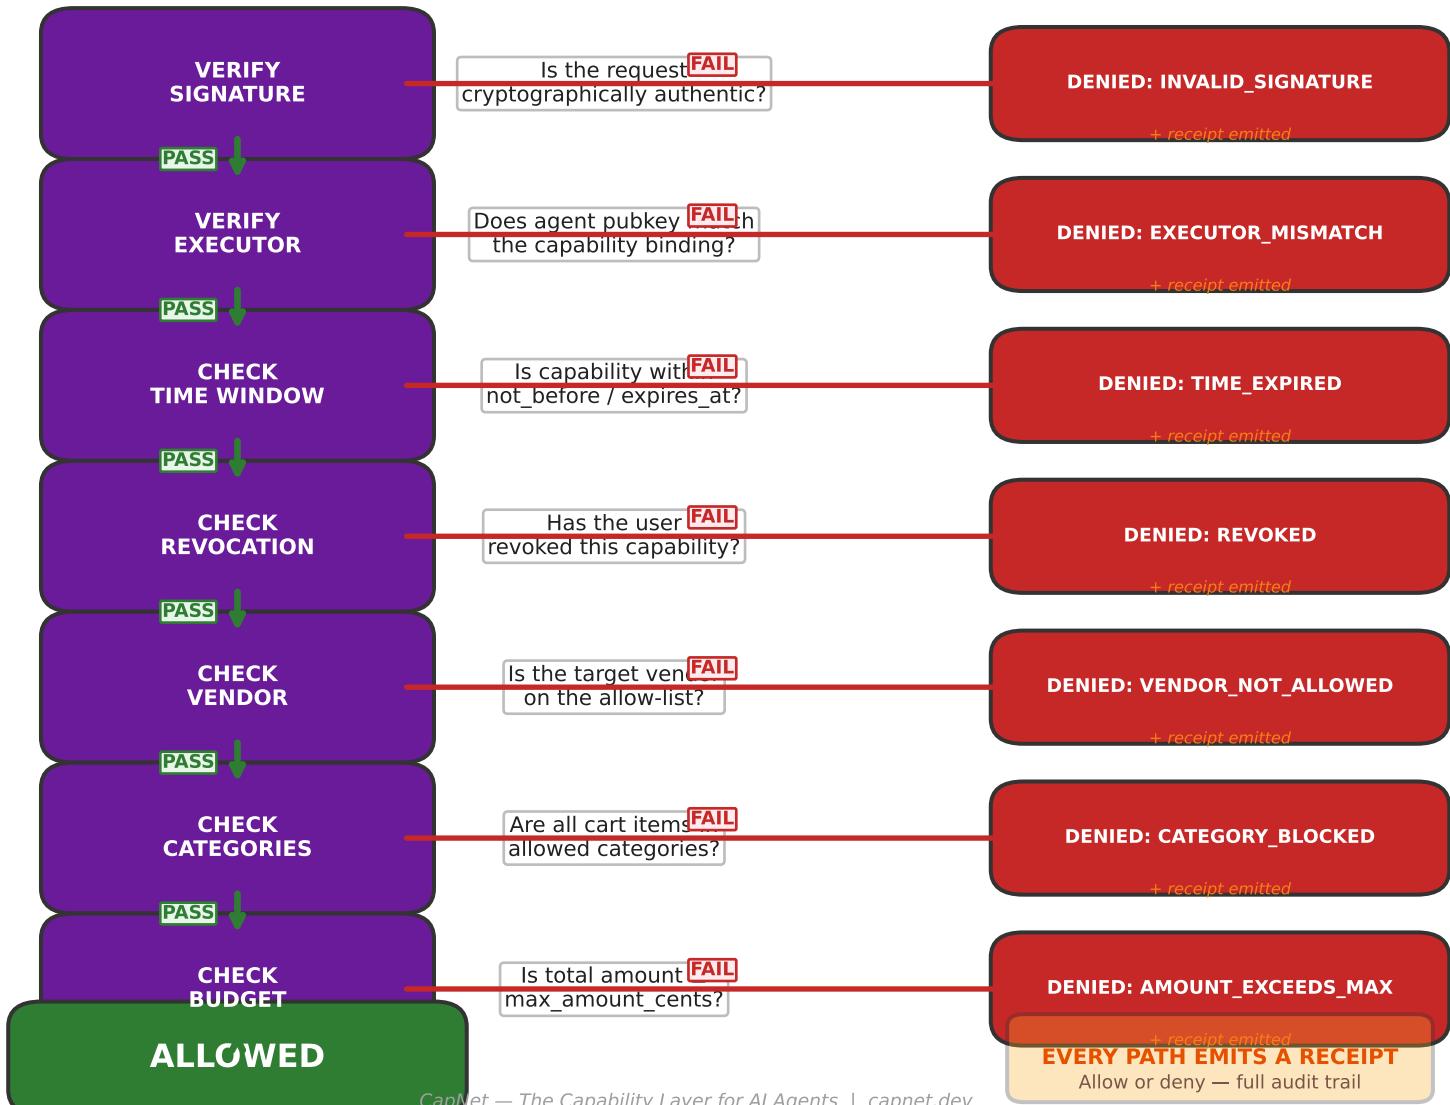
6

KEY INSIGHT: The agent receives authority (capability), NOT credentials.

Even if the agent is fully compromised, it cannot exceed the capability's constraints.

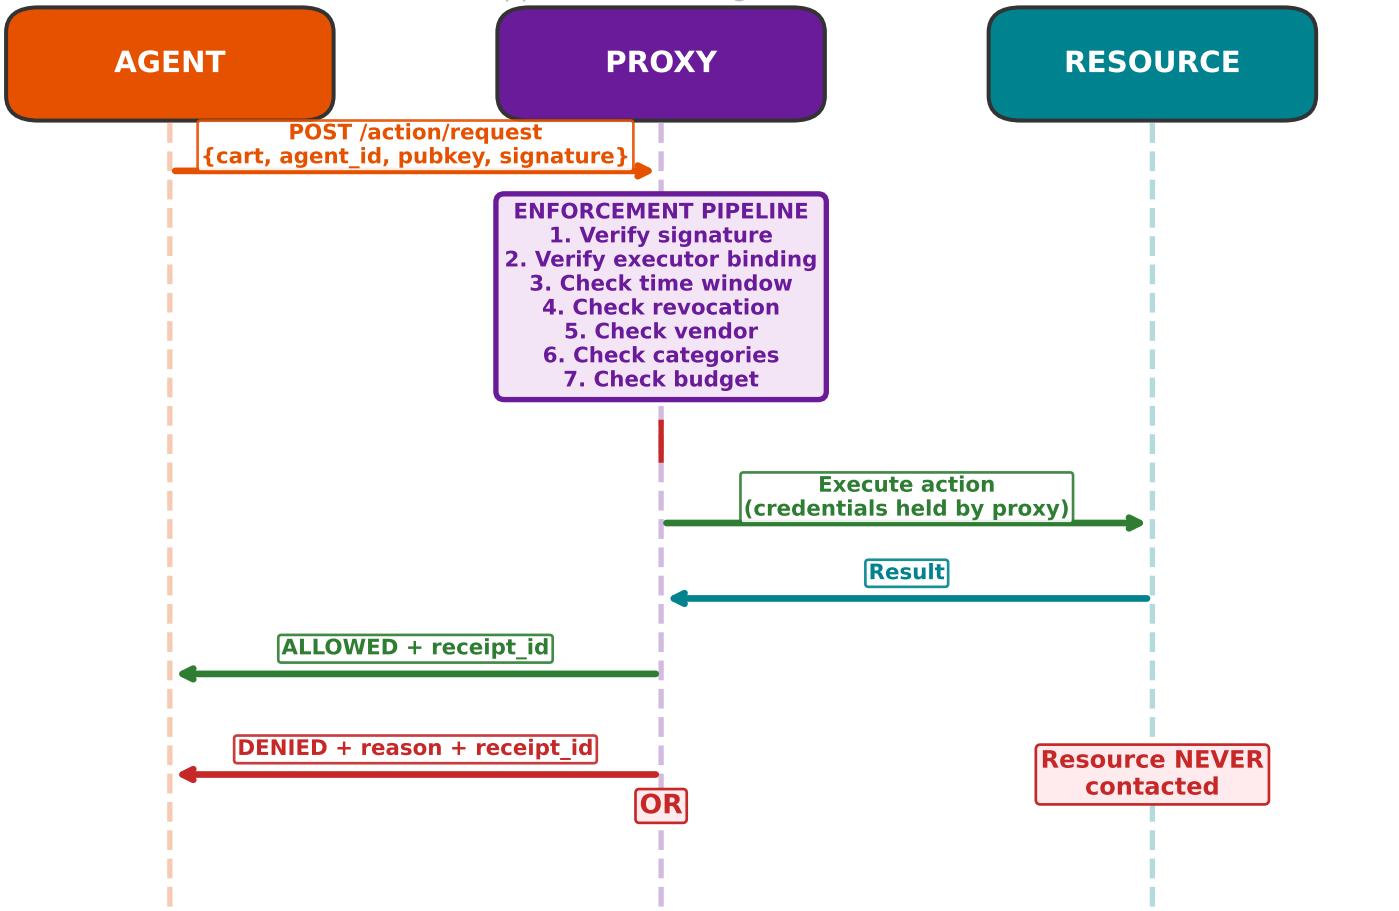
ENFORCEMENT DECISION TREE

Enforcement request passes through this pipeline — no exceptions



AGENT ACTION FLOW

What happens when an agent tries to take an action

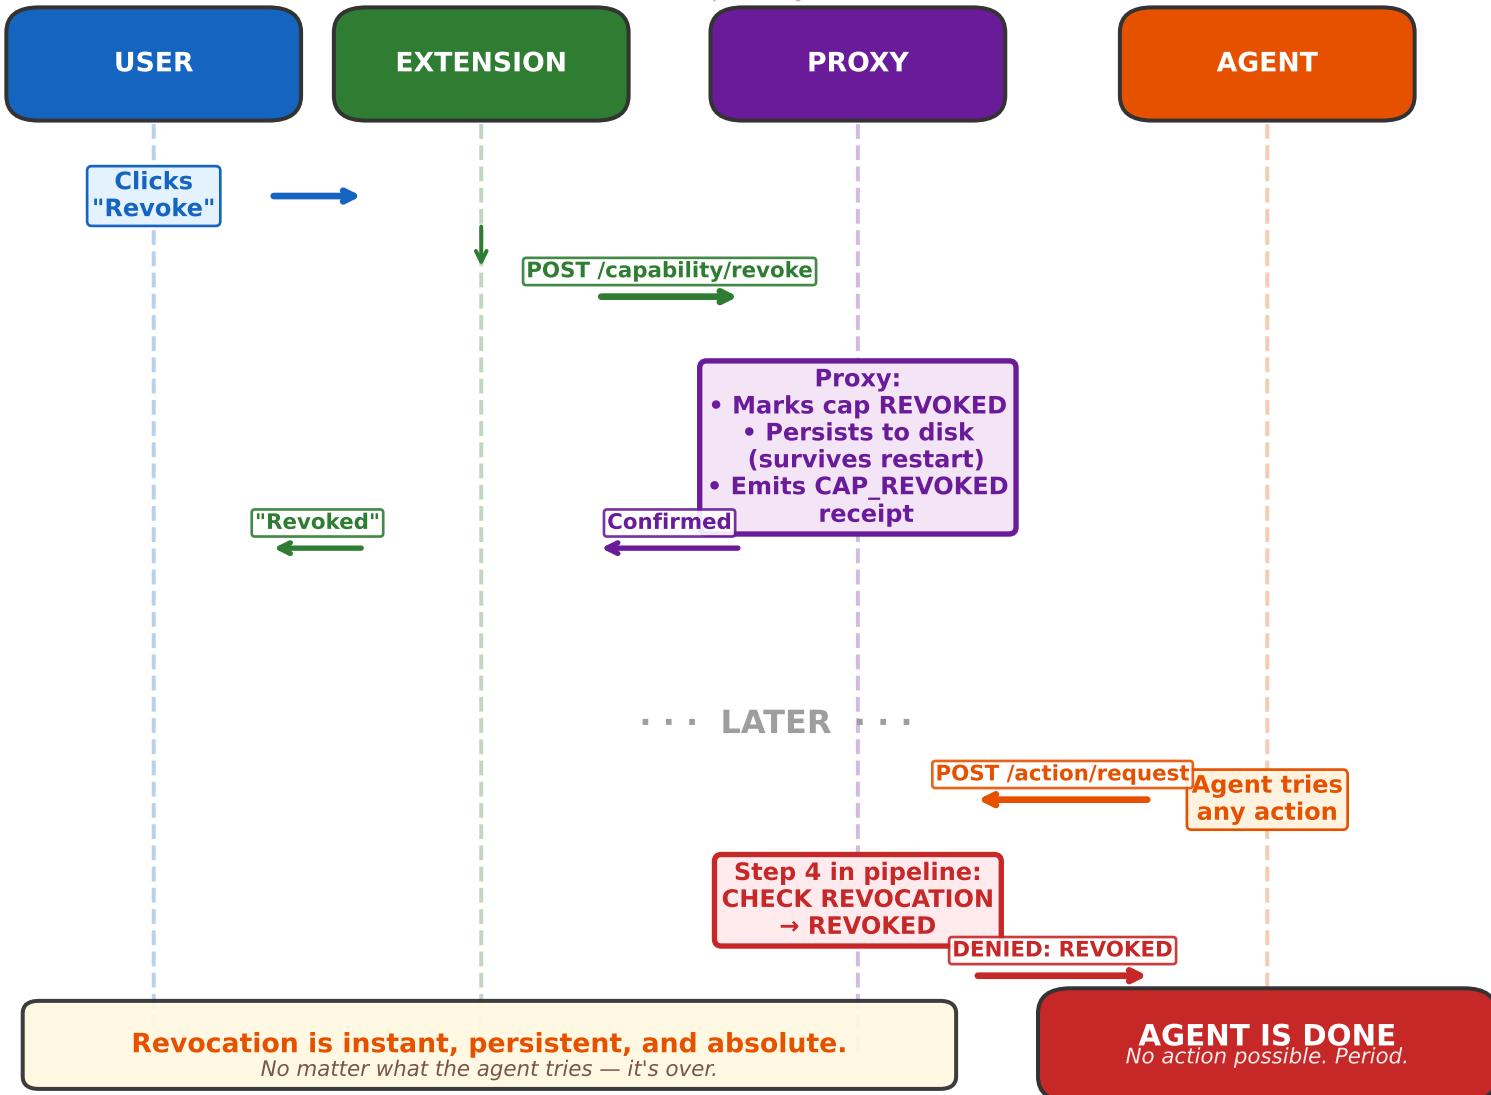


AUDIT TRAIL

Every request generates a signed receipt: ACTION_ATTEMPT → ACTION_ALLOWED or ACTION_DENIED
"Why did this happen?" is always answerable.

REVOCATION FLOW — KILL SWITCH

Instant capability termination



HIJACKER BLAST RADIUS

What happens when an agent is fully compromised

HIJACKER TAKES OVER AGENT

HAS ACCESS TO:

- ✓ Agent's Ed25519 keypair
- ✓ Knowledge of proxy API address
- ✓ Knowledge of capability ID

CAN DO:

- ✓ Send requests to proxy

CANNOT ACCESS:

- ✗ Merchant / service credentials
- ✗ Issuer signing key
- ✗ Other agents' keys
- ✗ Direct access to merchant API

CANNOT DO:

- ✗ Buy blocked categories
- ✗ Revocation controls
- ✗ Exceed budget limit
- ✗ Use unauthorized vendors
- ✗ Act after revocation
- ✗ Forge new capabilities

WORST CASE SCENARIO

✗ Escalate privileges

Hijacker can spend the remaining budget on allowed items at allowed vendors.

That's it. The blast radius IS the capability. User hits revoke → game over.

Compare: Traditional approach (shared credentials) → hijacker has FULL ACCESS to everything.

CAPNET vs TRADITIONAL APPROACHES

Why existing solutions don't solve the agent authorization problem

	API Keys / Credentials	OAuth Scopes	IAM / RBAC	CAPNET
Scoped authority	✗	~	~	✓
Time-bounded	✗	~	✗	✓
Instant revocation	✗	~	~	✓
Agent-specific binding	✗	✗	✗	✓
Budget enforcement	✗	✗	✗	✓
Category blocking	✗	✗	✗	✓
Vendor allow-listing	✗	✗	✗	✓
Delegation / attenuation	✗	✗	✗	✓
Audit trail (receipts)	✗	~	~	✓
Agent never sees creds	✗	✗	✗	✓
Survives agent compromise	✗	✗	✗	✓

✓ = Full support

~ = Partial / limited

✗ = Not supported

CapNet is purpose-built for the agent era.

OAuth answers "who is this?" — CapNet answers "what can this agent do right now, and can I stop it?"